


Bringing the Power of eBPF to Open vSwitch

Linux Plumber 2018

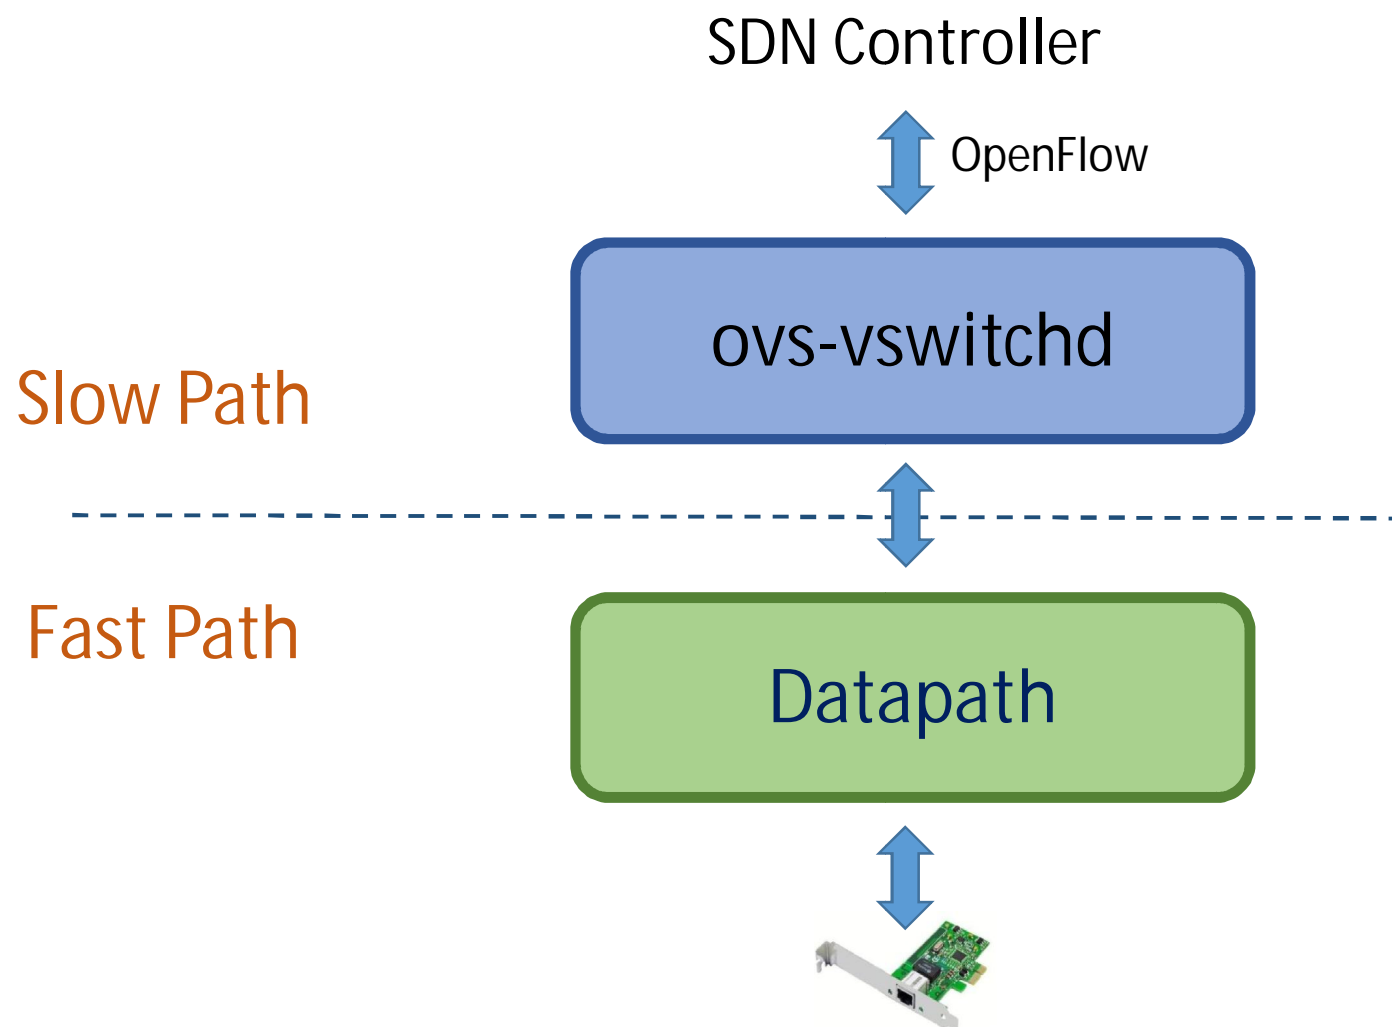
William Tu, Joe Stringer, Yifeng Sun, Yi-Hung Wei
VMware Inc. and Cilium.io



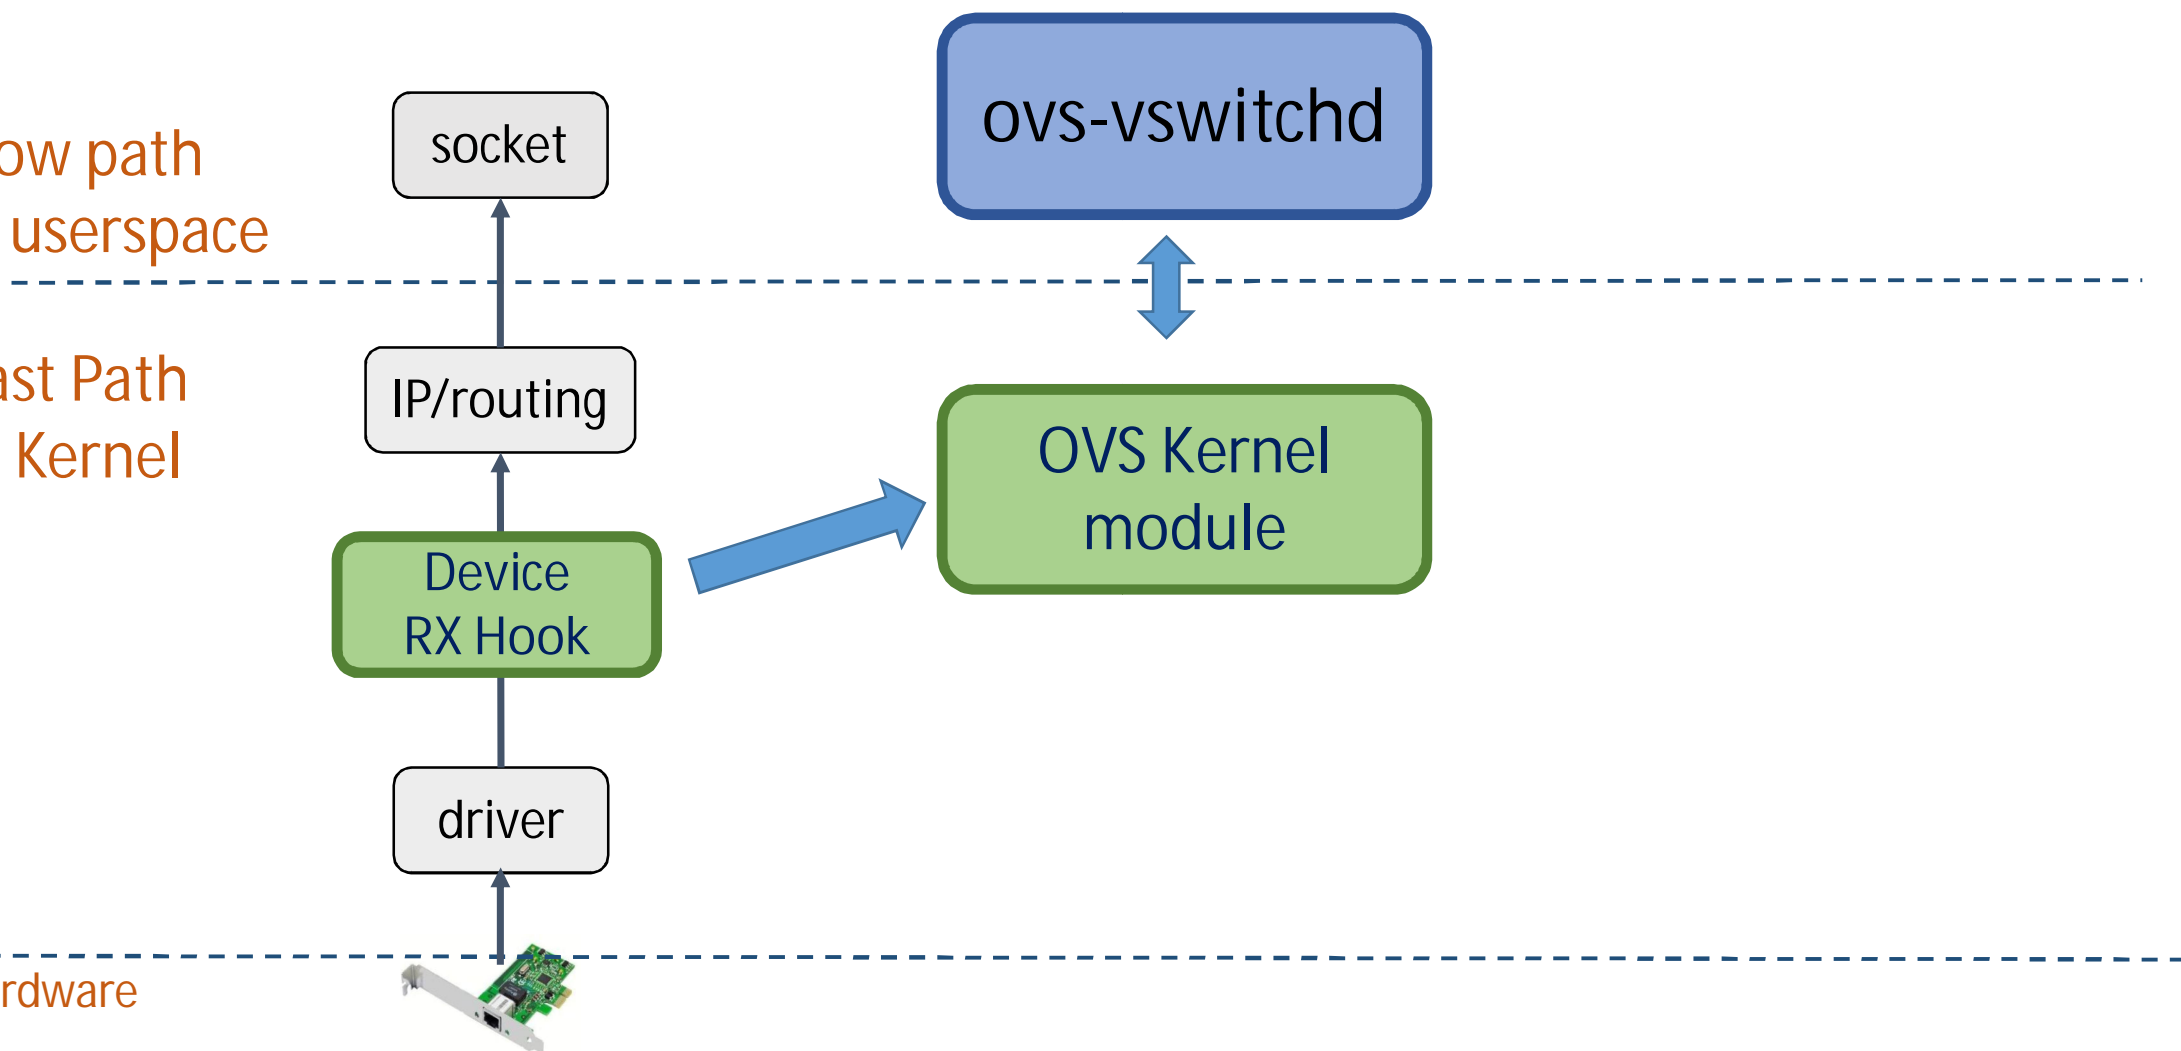
Outline

- Introduction and Motivation
- OVS-eBPF Project
- OVS-AF_XDP Project
- Conclusion

What is OVS?




OVS Linux Kernel Datapath



OVS-eBPF

OVS-eBPF Motivation

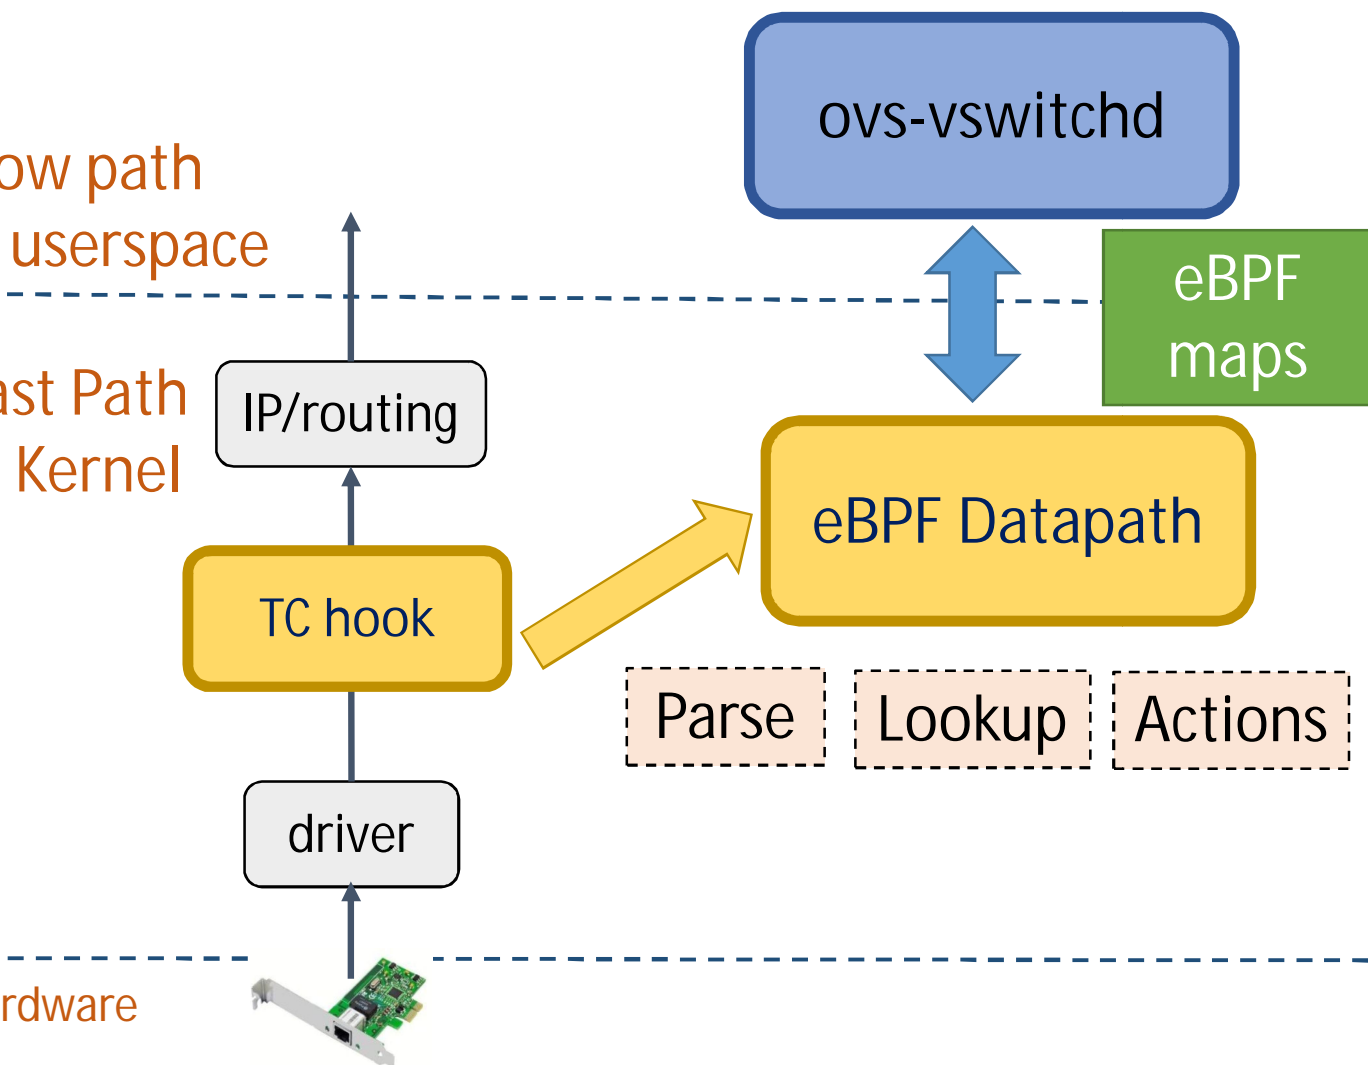
- Maintenance cost when adding a new datapath feature:
 - Time to upstream and time to backport
 - Maintain ABI compatibility between different kernel and OVS versions.
 - Different backported kernel, ex: RHEL, grsecurity patch
 - Bugs in compat code are easy to introduce and often non-obvious to fix
- Implement datapath functionalities in eBPF
 - Reduce dependencies on different kernel versions
 - More opportunities for experiments



What is eBPF?

- A way to write a **restricted C** program and runs in **Linux kernel**
 - A virtual machine running in Linux kernel
 - Safety guaranteed by BPF verifier
- Maps
 - Efficient key/value store resides in kernel space
 - Can be shared between eBPF program and user space applications
- Helper Functions
 - A core kernel defined set of functions for eBPF program to retrieve/push data from/to the kernel

OVS-eBPF Project



Goal

- Re-write OVS kernel datapath **entirely** with eBPF
- ovs-vswitchd controls and manages the eBPF DP
- eBPF map as channels in between
- eBPF DP will be specific to ovs-vswitchd

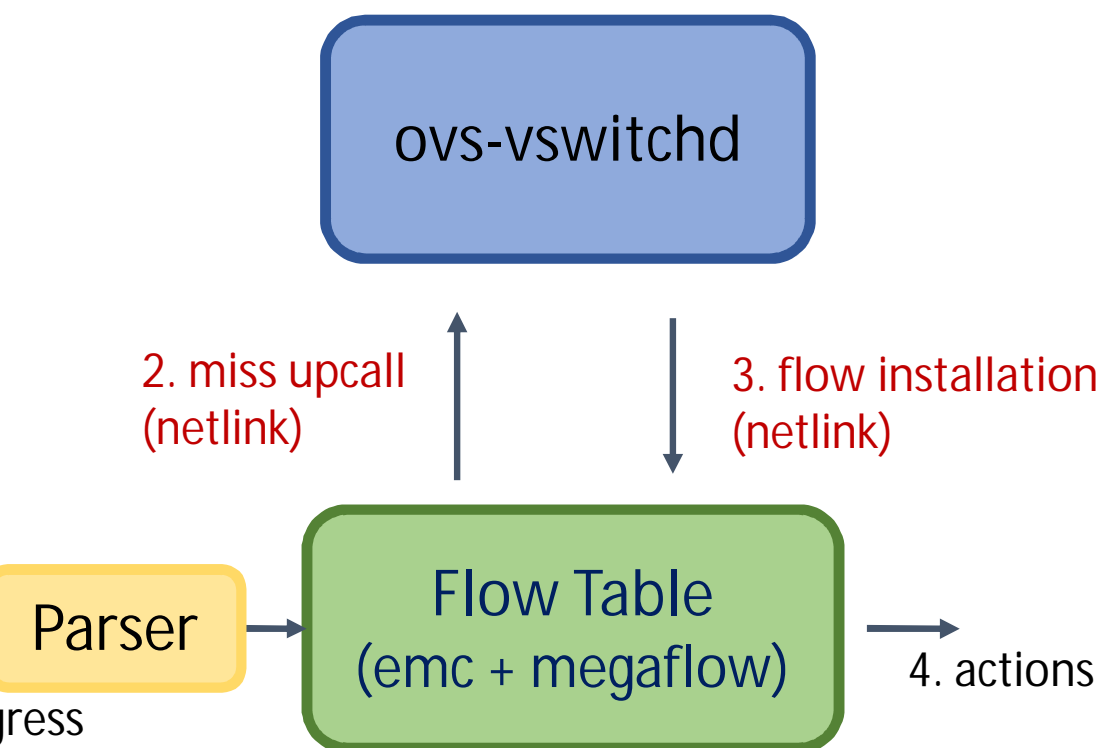
Headers/Metadata Parsing

- Define a flow key similar to struct `sw_flow_key` in kernel
- Parse protocols on packet data
- Parse metadata on struct `__sk_buff`
- Save flow key in per-cpu eBPF map

Difficulties

- Stack is heavily used
- Program is very branchy

Review: Flow Lookup in Kernel Datapath



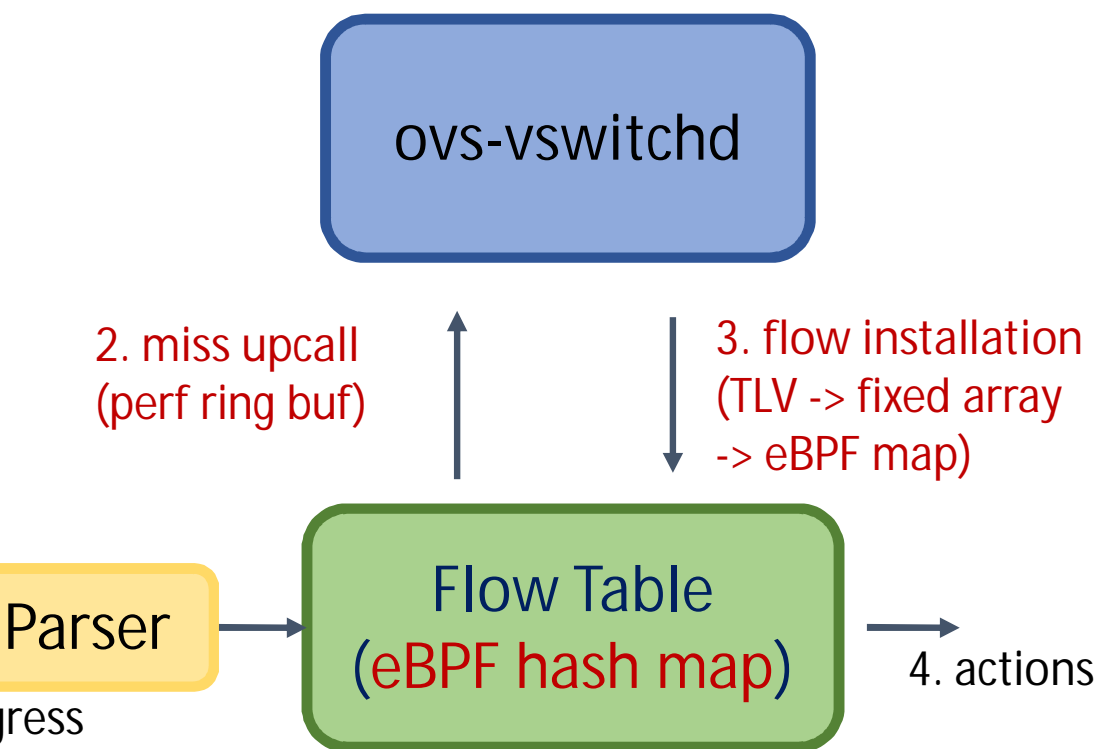
Slow Path

- Ingress: lookup miss and upcall
- ovs-vswitchd receives, does flow translation, and programs flow entry into flow table in OVS kernel module
- OVS kernel DP installs the flow entry
- OVS kernel DP receives and executes actions on the packet

Fast Path

- Subsequent packets hit the flow cache

Flow Lookup in eBPF Datapath



Slow Path

- Ingress: lookup miss and upcall
- Perf ring buffer carries packet and its metadata to ovs-vswitchd
- ovs-vswitchd receives, does flow translation, and programs flow entry into **eBPF map**
- ovs-vswitchd sends the packet down - trigger lookup again

Fast Path

- Subsequent packets hit the flow cache


Limitation on flow installation:

TLV format currently not supported in BPF verifier

Solution: Convert TLV into fixed length array

Review: OVS Kernel Datapath Actions

A list of actions to execute on the packet




Example cases of DP actions

- Flooding:
 - Datapath actions= output:9,output:5,output:10,...
- Mirror and push vlan:
 - Datapath actions= output:3,push_vlan(vid=17,pcp=0),output:2
- Tunnel:
 - Datapath actions:
 set(tunnel(tun_id=0x5,src=2.2.2.2,dst=1.1.1.1,ttl=64,flags(df|key))),output:1

eBPF Datapath Actions

A list of actions to execute on the packet



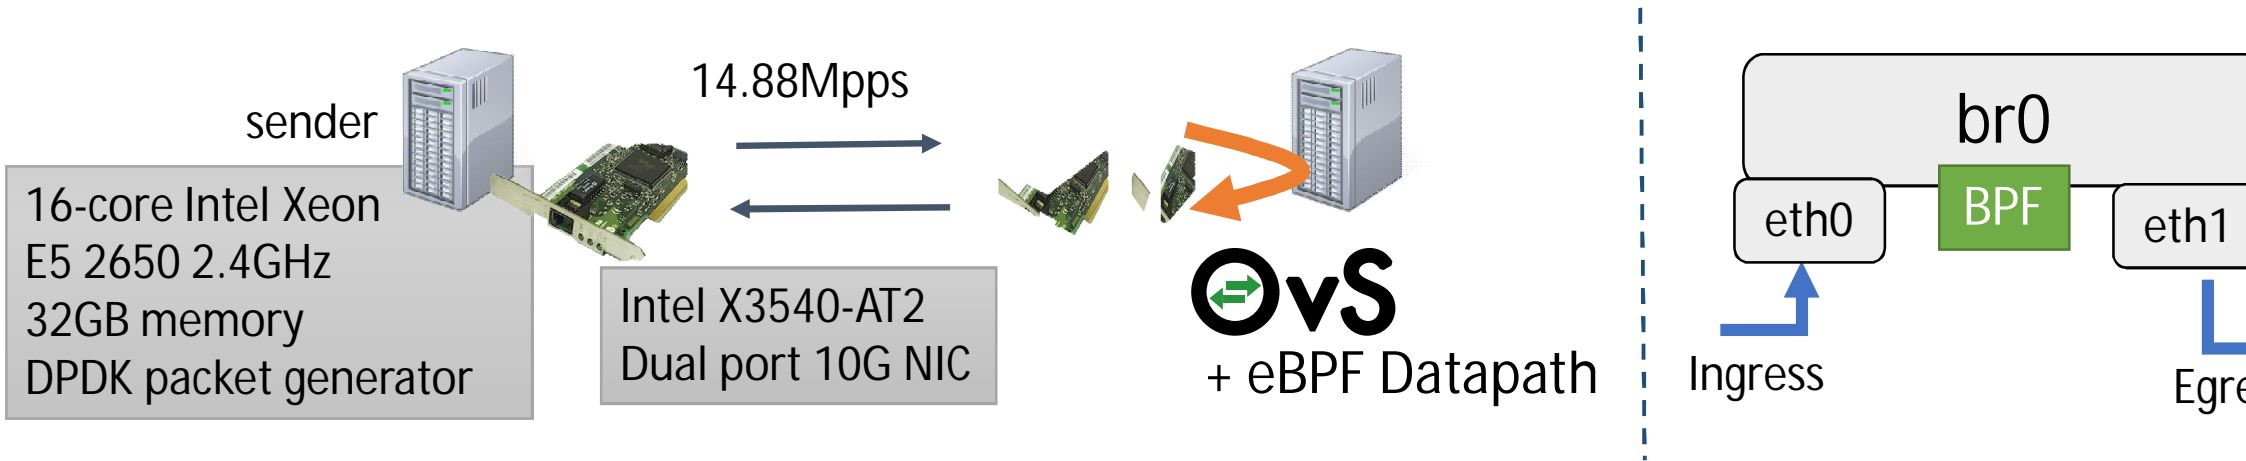
Challenges

- Limited eBPF program size (maximum 4K instructions)
- Variable number of actions: BPF disallows loops to ensure program termination

Solution:

- Make each action type an eBPF program, and tail call the next action
- Side effects: tail call has **limited context** and **does not return**
- Solution: keep **action metadata** and **action list** in a map

Performance Evaluation



- Sender sends 64Byte, 14.88Mpps to one port, measure the receiving packet rate at the other port
- OVS receives packets from one port, forwards to the other port
- Compare OVS kernel datapath and eBPF datapath
- Measure single flow, single core performance with Linux kernel 4.9-rc3 on OVS server

OVS Kernel and eBPF Datapath Performance

OVS Kernel DP Actions	Mpps
Output	1.34
Set dst_mac	1.23
Set GRE tunnel	0.57

eBPF DP Actions	Mpps
Redirect _(no parser, lookup, actions)	1.90
Output	1.12
Set dst_mac	1.14
Set GRE tunnel	0.48

All measurements are based on single flow, single core.

Conclusion and Future Work

Features

- Megaflow support and basic conntrack in progress
- Packet (de)fragmentation and ALG under discussion


Lesson Learned

- Writing large eBPF code is still hard for experienced C programmers
- Lack of debugging tools
- OVS datapath logic is difficult

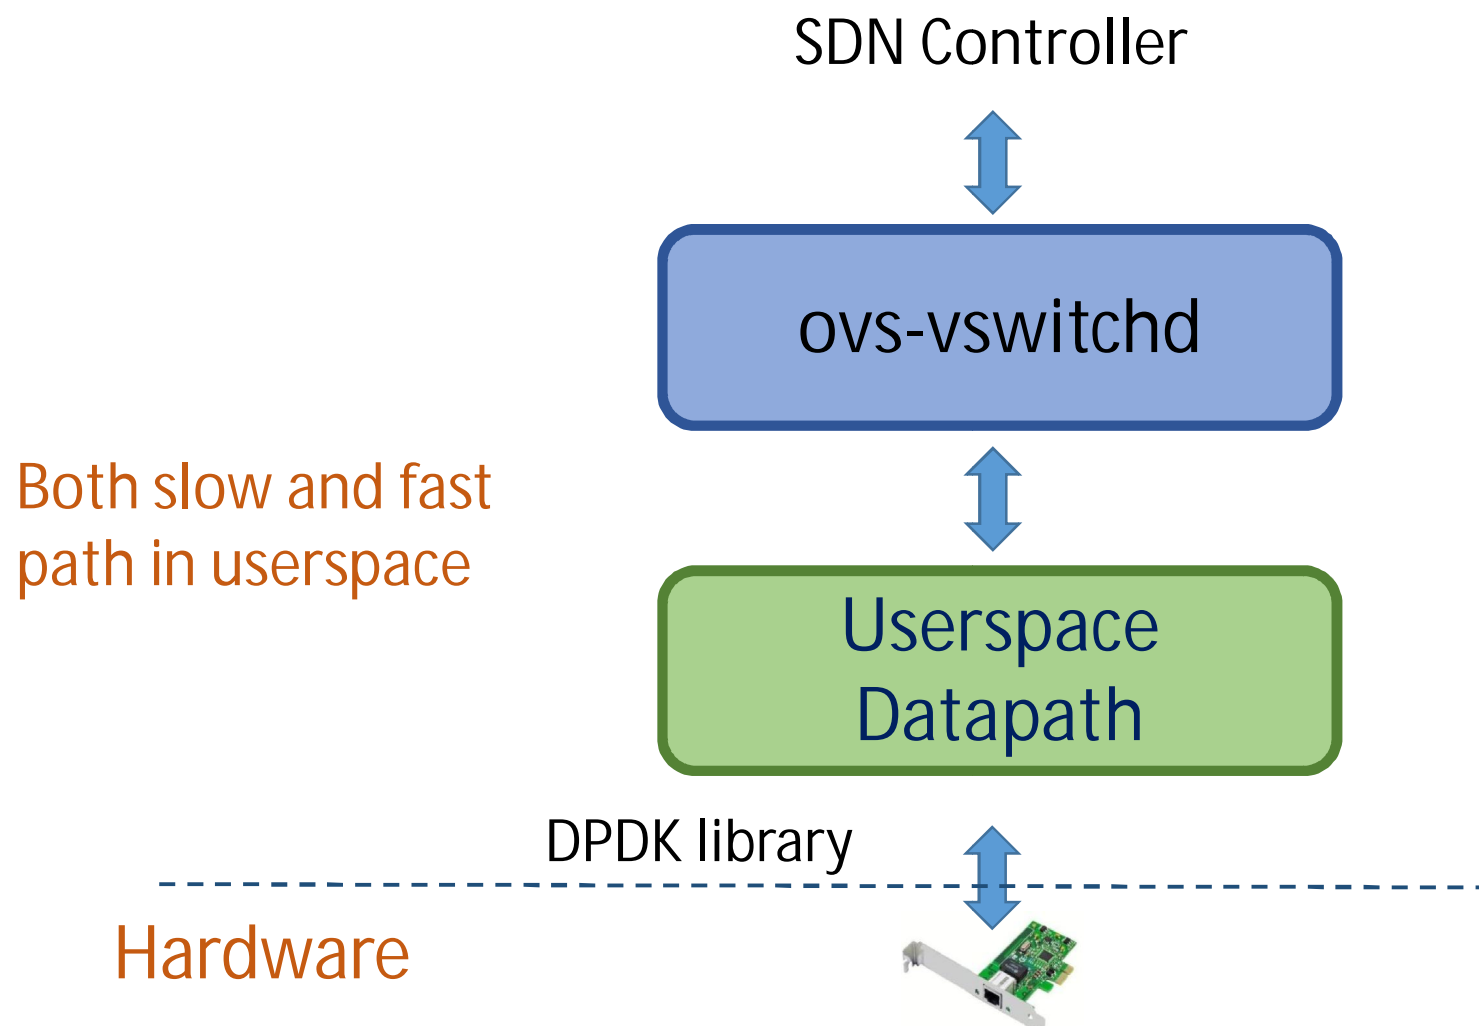
OVS-AF_XDP

OVS-AF_XDP Motivation

- Pushing all OVS datapath features into eBPF is hard
 - A large flow key on stack
 - Variety of protocols and actions
 - Dynamic number of actions applied for each flow
- Idea
 - Retrieve packets from kernel as fast as possible
 - Reuse the userspace datapath for flow processing
 - Less kernel compatibility than OVS kernel module

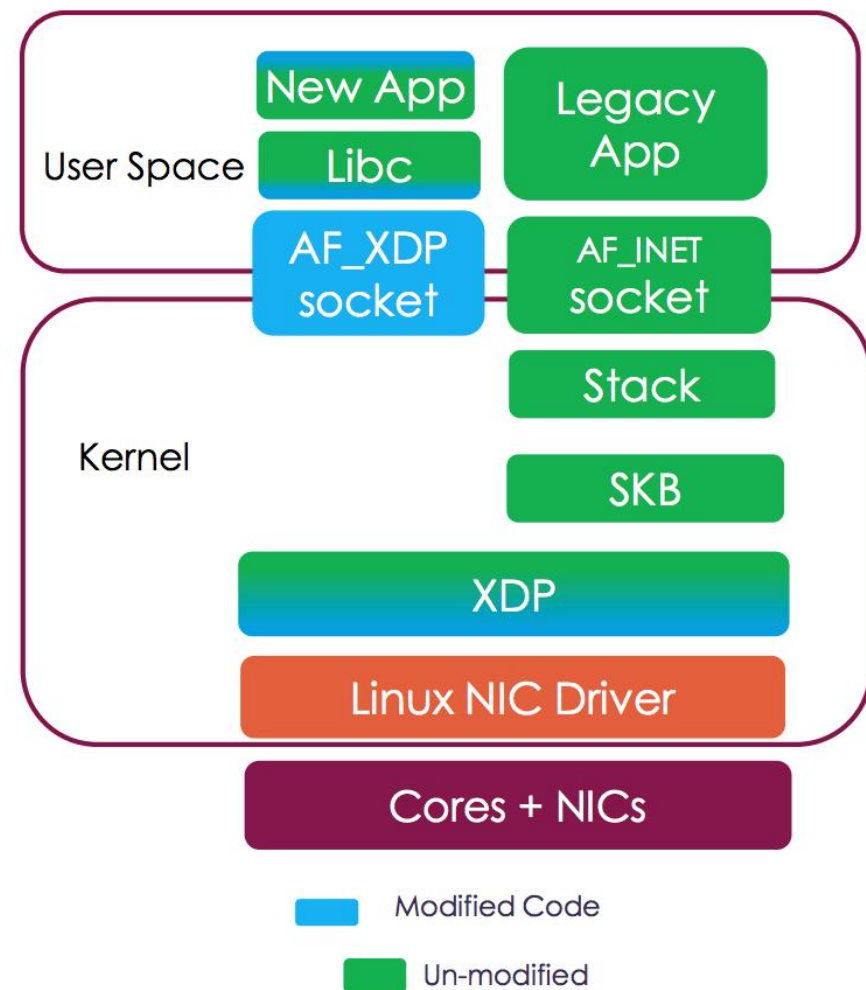


OVS Userspace Datapath (dpif-netdev)



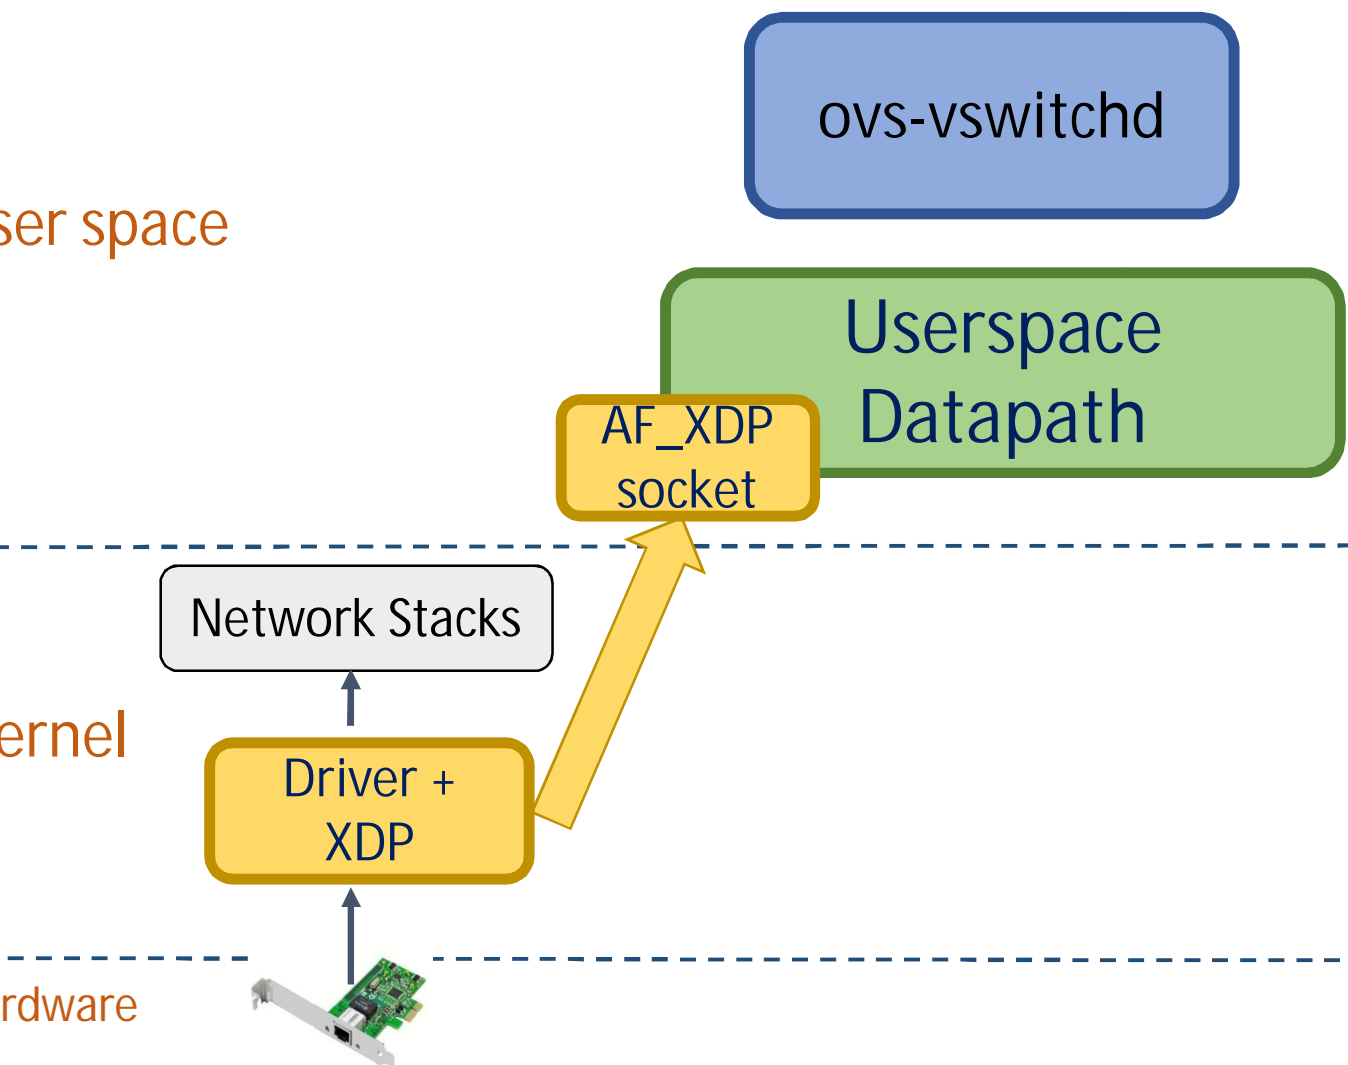
XDP and AF_XDP

- XDP: eXpress Data path
 - An eBPF hook point at the network device driver level
- AF_XDP:
 - A new **socket type** that receives/sends raw frames with high speed
 - Use XDP program to trigger receive
 - Userspace program manages Rx/Tx ring and Fill/Completion ring.
 - Zero Copy from DMA buffer to user space memory, **umem**



From "DPDK PMD for AF_XDP"

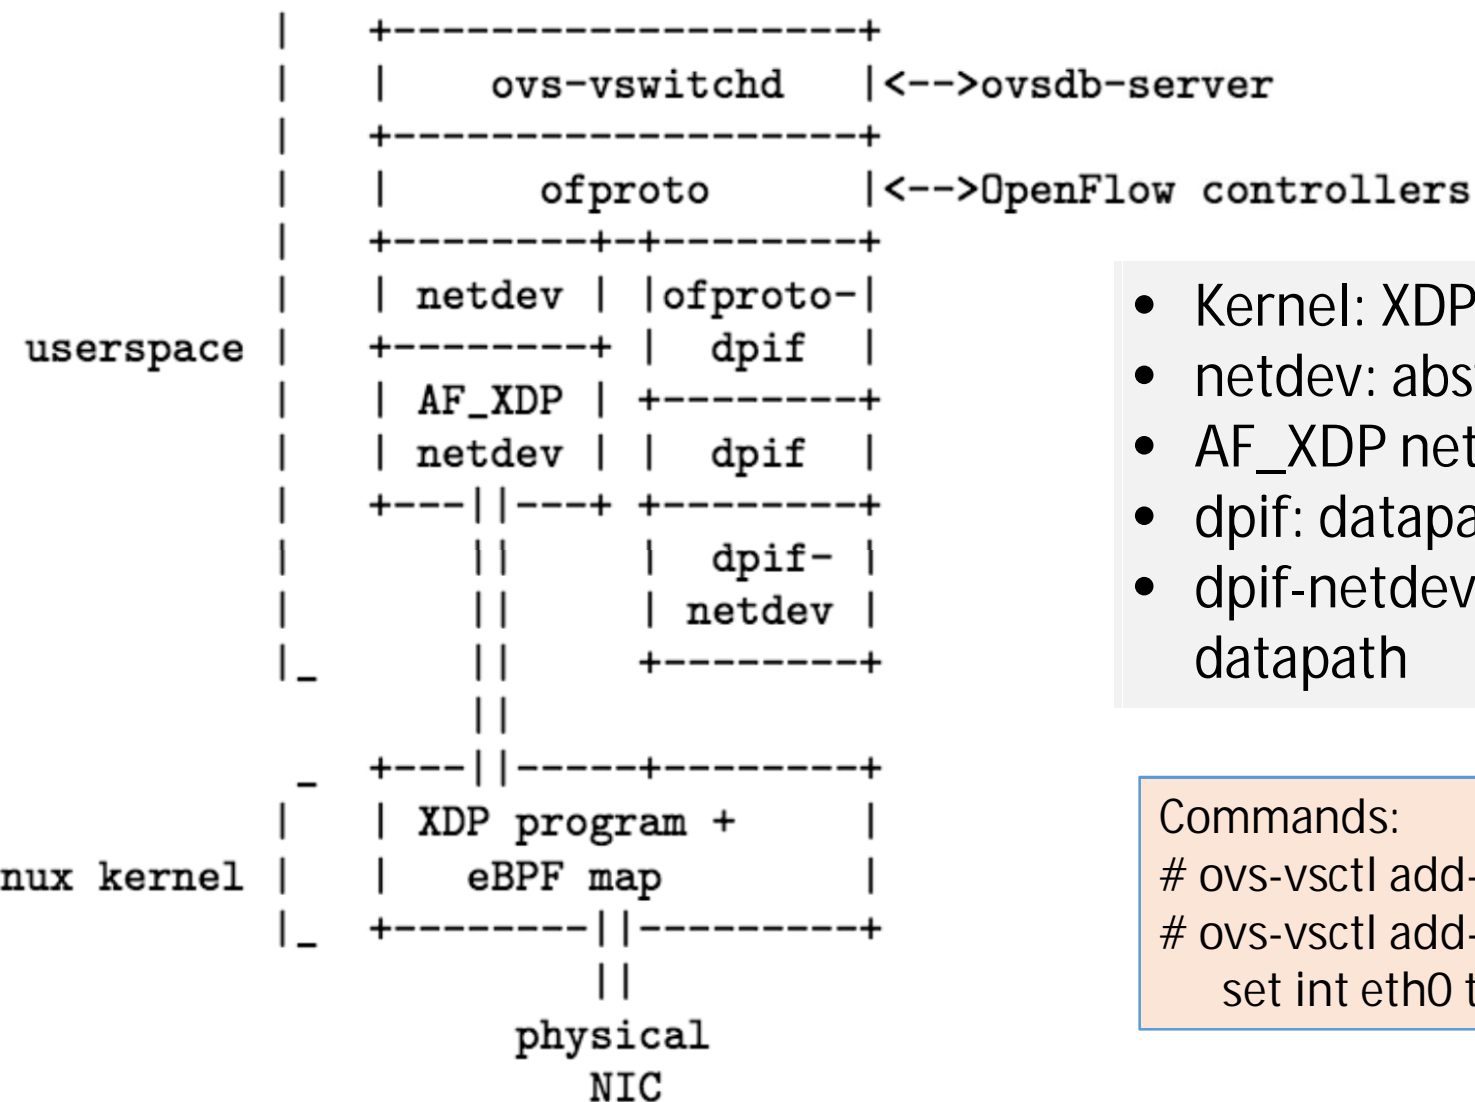
OVS-AF_XDP Project



Goal

- Use AF_XDP socket as a fast channel to userspace OVS datapath
- Flow processing happens in userspace

OVS-AF_XDP Architecture



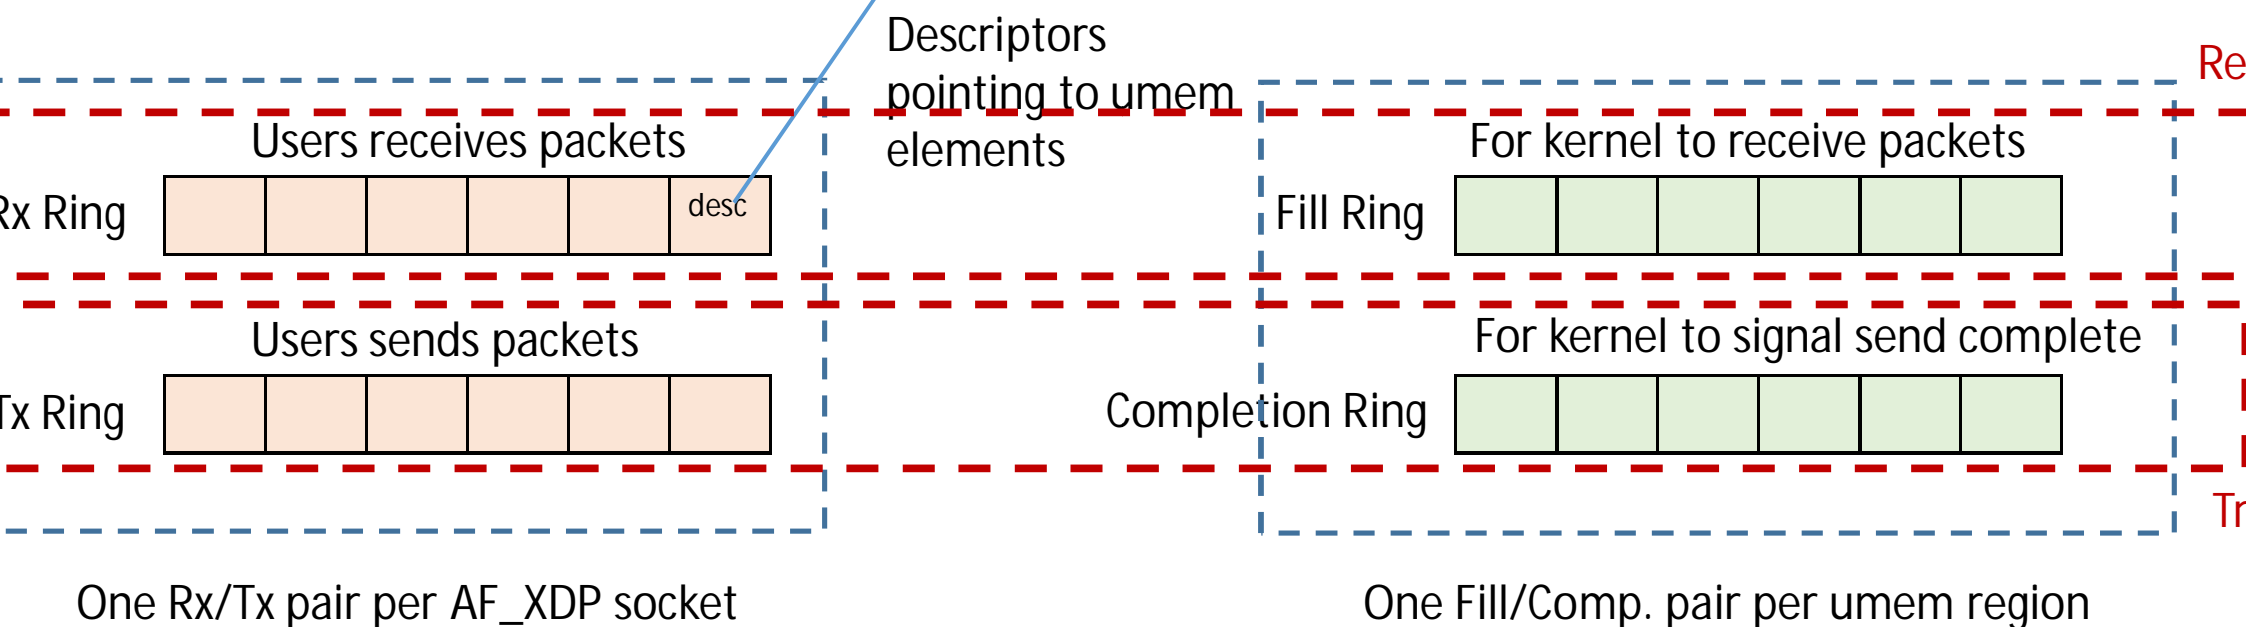

- Kernel: XDP program and eBPF map
- netdev: abstraction layer for network device
- AF_XDP netdev: implementation of afxdp device
- dpif: datapath interface
- dpif-netdev: userspace implementation of OVS datapath

Commands:

```
# ovs-vsctl add-br br0
# ovs-vsctl add-port br0 eth0 --
    set int eth0 type="afxdp"
```

AF_XDP umem and rings Introduction



umem memory region: multiple 2KB chunk elements



One Rx/Tx pair per AF_XDP socket


One Fill/Comp. pair per umem region

OVS-AF_XDP: Packet Reception (0)




OVS-AF_XDP: Packet Reception (1)

GET four elements,
program to Fill ring



OVS-AFXDP: Packet Reception (2)


umem consisting of 8 elements



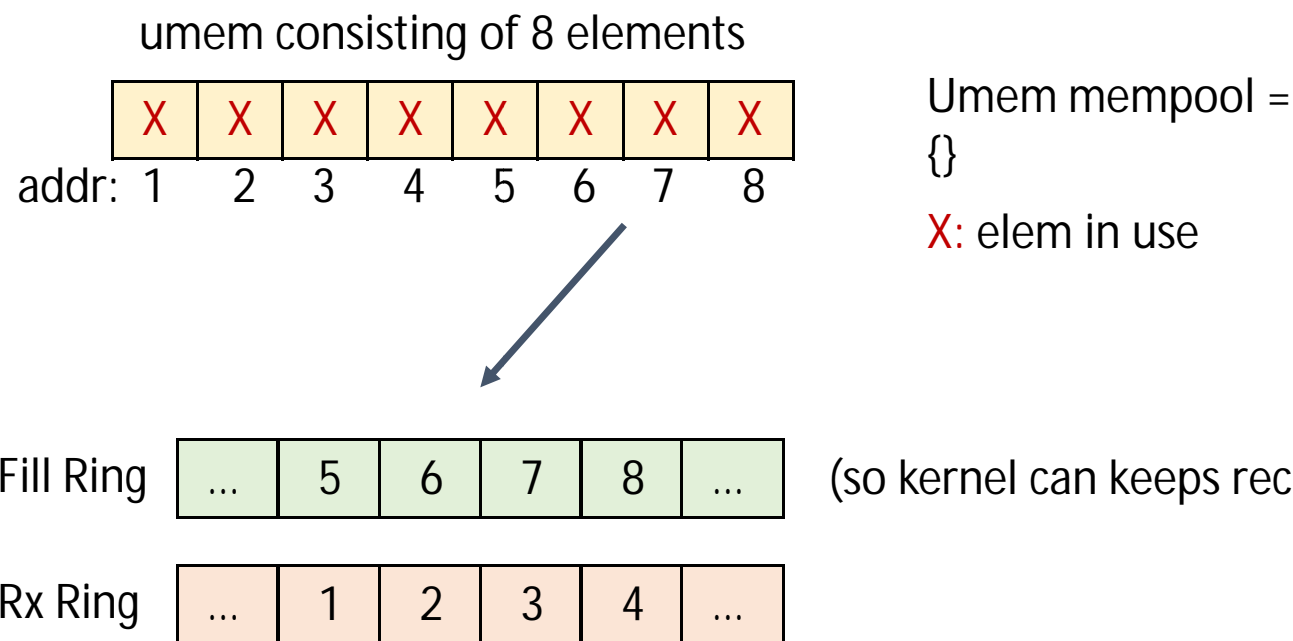
Umem mempool =
{5, 6, 7, 8}

X: elem in use

Kernel receives four packets
 puts them into the four umem chunks
 transition to Rx ring for users




OVS-AFXDP: Packet Reception (3)



OVS-AFXDP: Packet Reception (4)

umem consisting of 8 elements




Umem mempool =
{

X: elem in use


OVS userspace processes packets
on Rx ring



OVS-AFXDP: Packet Reception (5)




OVS-AFXDP: Packet Transmission (0)



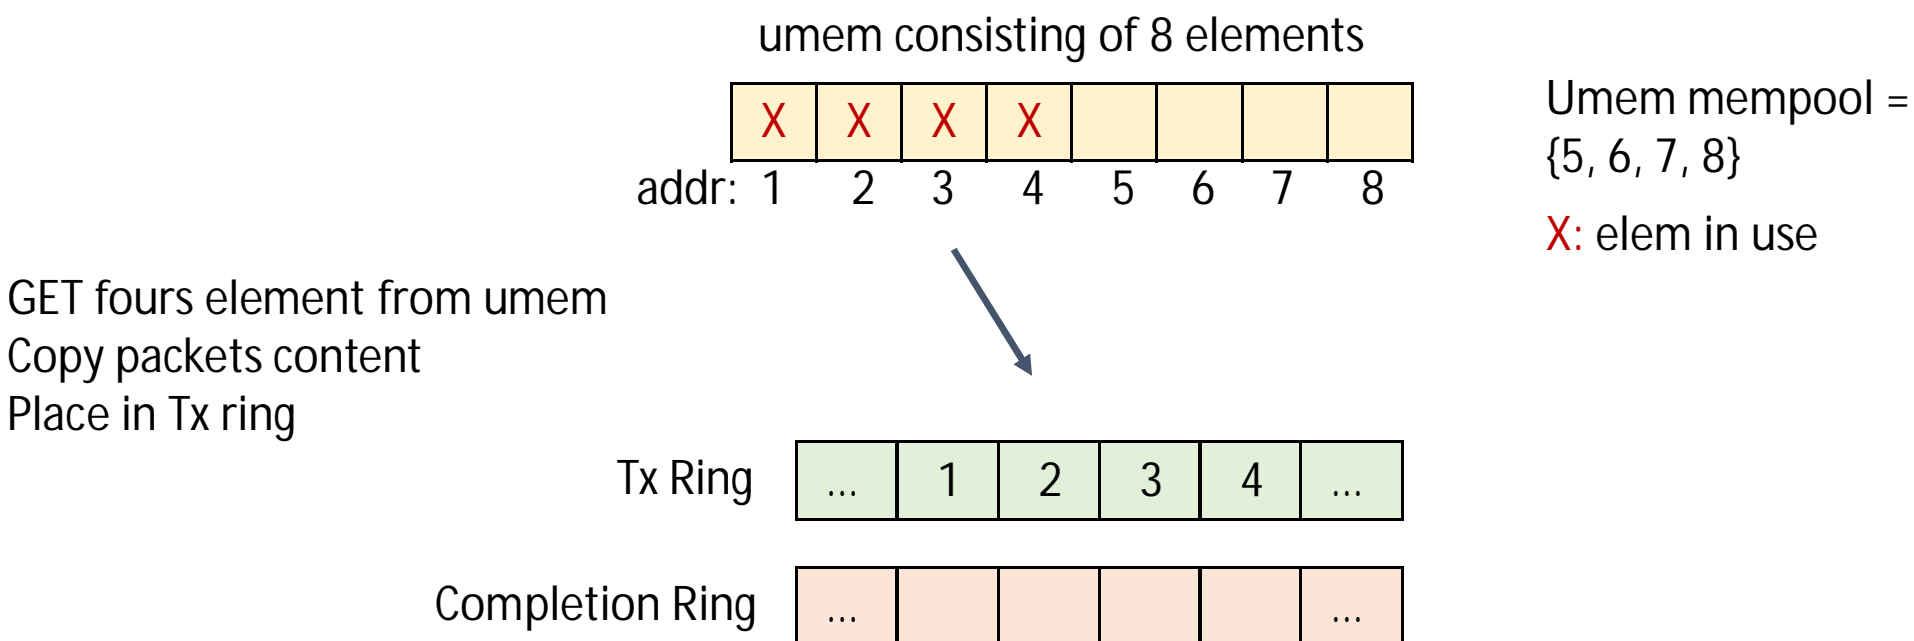
Umem mempool =
{1, 2, 3, 4, 5, 6, 7, 8}

X: elem in use

OVS userspace has four packets to send




OVS-AFXDP: Packet Transmission (1)



OVS-AFXDP: Packet Transmission (2)

umem consisting of 8 elements



Umem mempool =
{5, 6, 7, 8}


X: elem in use

Issue sendmsg() syscall
Kernel tries to send packets
on Tx ring



OVS-AFXDP: Packet Transmission (3)


umem consisting of 8 elements




Umem mempool =
{5, 6, 7, 8}

X: elem in use

Kernel finishes sending
Transition the four elements
to Completion Ring for users





OVS-AFXDP: Packet Transmission (4)



Umem mempool = {1, 2, 3, 4, 5, 6, 7, 8}

X: elem in use

OVS knows send operation is done
Recycle/PUT the four elements back
to umempool



Optimizations

- OVS pmd (Poll-Mode Driver) netdev for rx/tx
 - Before: call poll() syscall and wait for new I/O
 - After: dedicated thread to busy polling the Rx ring
- UMEM memory pool
 - Fast data structure to GET and PUT umem elements
- Packet metadata allocation
 - **Before:** allocate md when receives packets
 - **After:** pre-allocate md and initialize it
- Batching sendmsg system call


Umempool Design

- umempool keeps track of available umem elements
 - GET: take out N umem elements
 - PUT: put back N umem elements
- Every ring access need to call umem element GET/PUT

Three designs:

- LIFO-List_head: embed in umem buffer, linked by a list_head, push/pop style
- FIFO-ptr_ring: a pointer ring with head and tail pointer
- LIFO-ptr_array: a pointer array and push/pop style access


LIFO-list_head Design



Idea:

- Reserve first 8-byte in umem buffer as next pointer
- No extra mempool metadata allocation needed
- GET from the head and PUT from the head


FIFO-ptr_ring Design



Idea:

- Allocate a ring with each ring element contains umem address
- Producer: PUT elements at `ptr_ring[head]` and `head++`
- Consumer: GET elements and `ptr_ring[tail]` and `tail++`

LIFO-ptr_array Design



Multiple 2K umem chunk memory region

Idea:


- Each ptr_array element contains a umem address
- Producer: PUT elements on top and top++
- Consumer: GET elements from top and top--

Packet Metadata Allocation

- Every packets in OVS needs metadata: struct dp_packet
- Initialize the packet data independent fields

Two designs:

1. Embedding in umem packet buffer:
 - Reserve first 256-byte for struct dp_packet
 - Similar to DPDK mbuf design
2. Separate from umem packet buffer:
 - Allocate an array of struct dp_packet
 - Similar to skb_array design

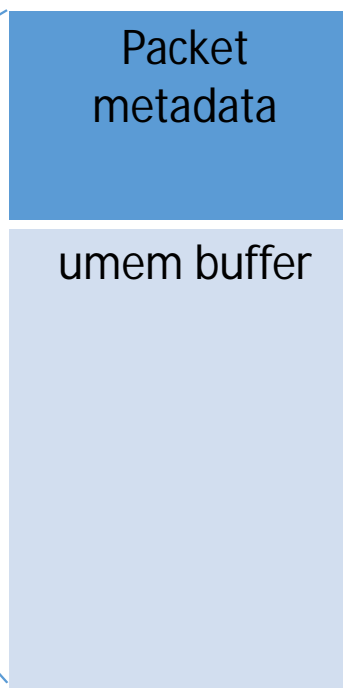



Packet Metadata Allocation

Embedding in umem packet buffer

Each 2K has:

Multiple 2K umem chunk memory region




Packet Metadata Allocation

Separate from umem packet buffer

Packet metadata in another memory region

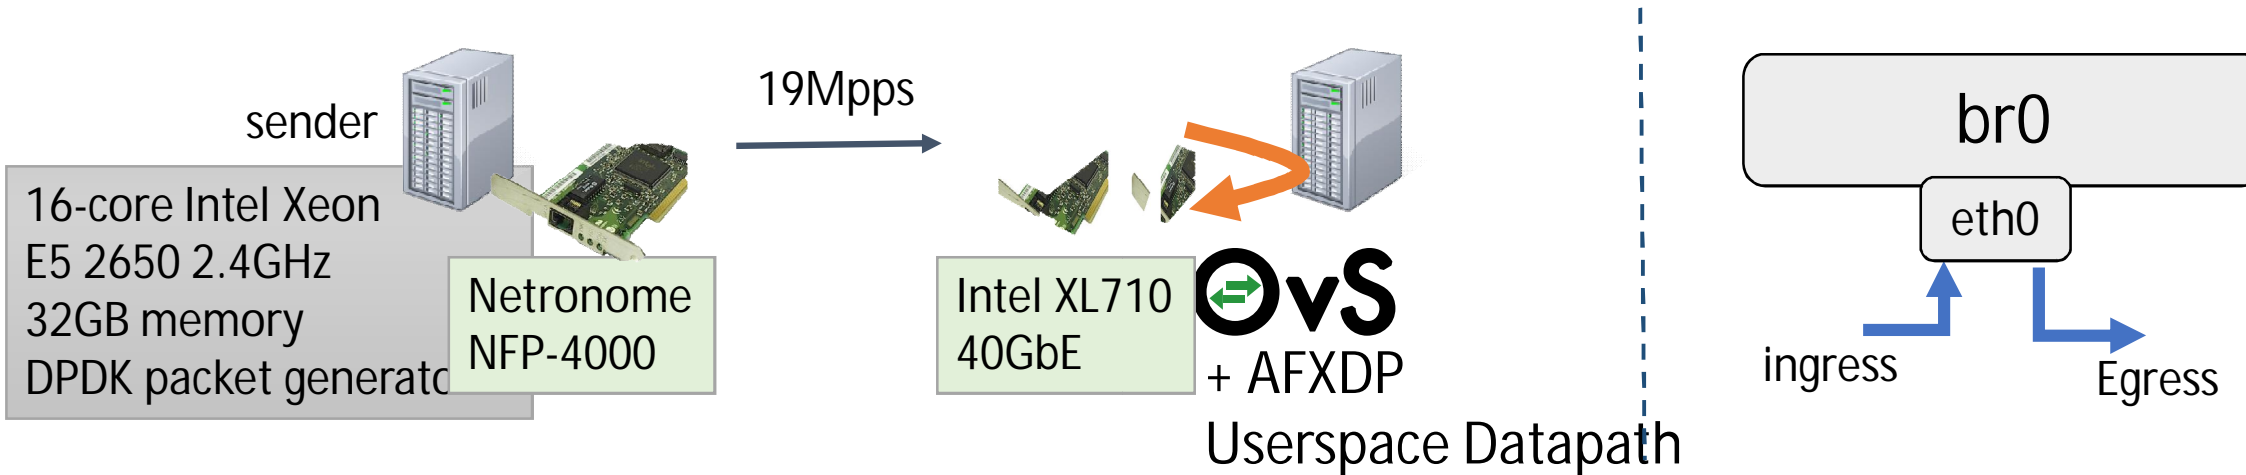


One-to-one maps to umem



Multiple 2K umem chunk memory region

Performance Evaluation



- Sender sends 64Byte, 19Mpps to one port, measure the receiving packet rate at the other port
- Measure single flow, single core performance with Linux kernel 4.19-rc3 and OVS 2.9
- Enable AF_XDP Zero Copy mode

Performance Evaluation

Experiments

- OVS-AFXDP
 - rxdrop: parse, lookup, and action = drop
 - l2fwd: parse, lookup, and action = set_mac, output to the received port
- XDPSOCK: AF_XDP benchmark tool
 - rxdrop/l2fwd: simply drop/fwd without touching packets
- LIFO-ptr_array + separate md allocation shows the best

Results

	XDPSOCK	OVS-AFXDP
rxdrop	19Mpps	19Mpps
l2fwd	17Mpps	14Mpps

Conclusion and Discussion

Future Work

- Follow up new kernel AF_XDP's optimizations
- Try virtual devices vhost/virtio with VM-to-VM traffic
- Bring feature parity between userspace and kernel datapath

Discussion

- Usage model: # of XSK, # of queue, # of pmd/non-pmd
- Comparison with DPDK in terms of deployment difficulty

Dislike?



Like?



Question?

Thank You

Batching sendmsg syscall

- Place a batch of 32 packets on TX ring, issue send syscall
- Design 1
- Check this 32 packets on completion ring, then recycle
- If not, keep issuing send
- Design 2
- Check any 32 packets on completion ring, then recycle
- If not, keep issuing send
- Design 3
- Issue sendmsg syscall